

Introduction to Statistical Methods

Augustin Luna
18 January, 2016

Research Fellow
Department of Biostatistics and Computational Biology
Dana-Farber Cancer Institute

Topics to be Covered

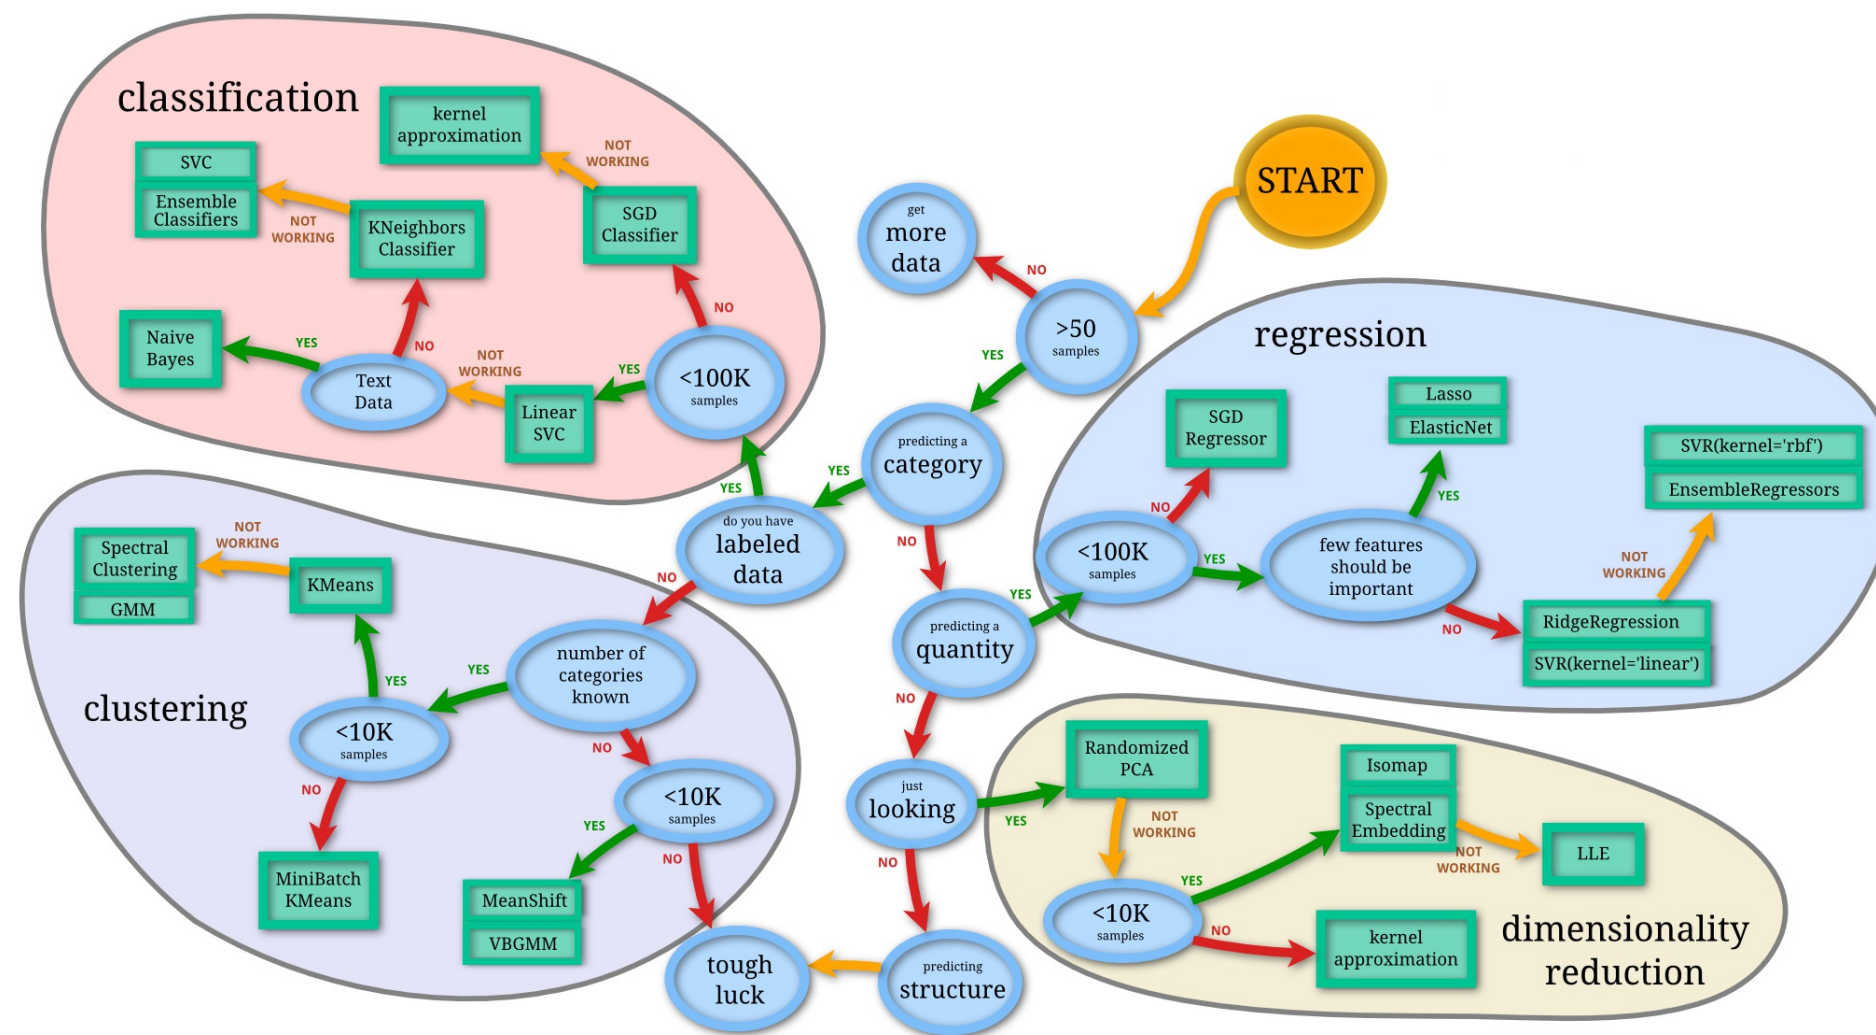
- Correlations
- Enrichment Analyses
- Multiple Testing Correction
- Regression
- Clustering
- Dimensionality Reduction

Introduction

- Data analyses are the product of many tasks
- Statistical Methods
 - Build predictive mathematical models
- Data preparation
 - Extracting structured data from unstructured data sources
 - Merging data sources
 - Ensuring consistency of datasets
- Dataset interpretation
 - Create visualizations to present and communicate findings
- Methods are common in the areas of informatics, data

Statistical Methods Flowchart

- The flowchart below helps find the right method for a given problem
 - http://scikit-learn.org/stable/tutorial/machine_learning_map/



Methods to be Covered

- Basic method will be covered to build confidence with R and the general concepts
 - Dimensionality Reduction: Principal Component Analysis (PCA)
 - Regression: Linear regression
 - Clustering: Hierarchical and K-means clustering
 - Classification will not be covered

Correlating Two Vectors

```
# Make sure the random numbers are always the same
set.seed(1)

# Generate two sets of 20 random numbers
a <- runif(20); b <- runif(20)

# Calculate the correlation of the two sets
cor.test(a, b)
```

Pearson's product-moment correlation

```
data:  a and b
t = -1.0368, df = 18, p-value = 0.3136
alternative hypothesis: true correlation is not equal to 0
95 percent confidence interval:
 -0.6152730  0.2292138
sample estimates:
      cor
-0.2373854
```

Extracting Values From Results

- Values in results are described in the help `?cor.test`
- A p-value is the probability of seeing results as extreme as the ones produced in an analysis.

```
set.seed(1)

a <- runif(20)
b <- runif(20)

results <- cor.test(a, b, method="pearson")
names(results)
```

```
[1] "statistic"      "parameter"      "p.value"        "estimate"
"null.value"
[6] "alternative"    "method"         "data.name"      "conf.int"
```

```
results$p.value
```


```
[1] 0.3135682
```

Over-Representation (ORA) and Enrichment Analyses

- Enrichment tests are widely used in biology to determine if the genes contain a trait more frequently than a random sampling of genes
 - Gene Ontology (GO) term (e.g. biological process, molecular function, or cellular component) and pathways are the most common comparisons made
- Several tools exist for doing enrichment analyses
 - The tests are either Fisher's exact test or hypergeometric test; these tests produce the same results
 - These calculations can be done in R using `fisher.test` and `phyper`

Enrichment Analysis

- What is the probability of randomly drawing at least 4 black points in a random sample of 5 points?
 - The concept of “black” could be replaced by “genes from a given pathway” or “genes with a common function”



Calculating an ORA (Enrichment) P-Value

The significance of an over-representation (enrichment analysis) is calculated using a hypergeometric test:

$$P(\text{pathway}_j) = 1 - \sum_{i=0}^{k-1} \frac{\binom{K}{i} \binom{N-K}{n-i}}{\binom{N}{n}}$$

where

- N: number of studied genes,
- n: total number of genes identified by a previous analysis
- K: total number of genes in a pathway pathway_j
- k: number genes previously identified genes in pathway_j
- $\binom{n}{k}$: choose k elements from a set with n elements disregarding order

The p-value of this test indicates the probability that a random selection of genes of the same size as the input gene set from a population would produce the same number of observed annotations (e.g. for a specific GO term or pathway) or more in the gene set

Choose k from N Without Order

```
N <- 5  
k <- 3  
factorial(N) / (factorial(k)*factorial(N-  
k))
```

```
[1] 10
```

```
N <- 5  
k <- 3  
choose(N, k)
```

```
[1] 10
```

Contingency Table for Enrichment Analysis

	Drawn	Not Drawn	Total
Black	$k=4$	$K-k=1$	$K=5$
Red	$n-k=6$	$N+k-n-K=39$	$N-K=45$
Total	$n=10$	$N-n=40$	$N=50$

phyper or fisher.test Example

- NOTE: `hitInSample-1` is necessary in `phyper` because if `lower.tail` is `FALSE`, probabilities returned are $P(X > x)$. Subtract `x` by 1 to get $P(X \geq x)$ (`x` equal to or greater than).

```
sampleSize <- 10 # size drawn
hitInSample <- 4 # black drawn
hitInPop <- 5 # all black
failInPop <- 50-hitInPop # number of red

phyper(hitInSample-1, hitInPop, failInPop, sampleSize,
lower.tail= FALSE);
```

```
[1] 0.004083521
```

```
fisher.test(matrix(c(hitInSample, hitInPop-hitInSample,
sampleSize-hitInSample, failInPop-sampleSize
+hitInSample), 2, 2), alternative='greater')$p.value;
```

```
[1] 0.004083521
```

Multiple Testing Correction for Enrichment Analyses

- Enrichment analyses often do analyses over a large number of molecular functions or pathways
- If we conduct many such tests, we are likely to see false positives
 - A p-value significance cutoff of 0.05 means that we expect 1 test out of 20 to appear significant by random chance (i.e. a false positive)

Multiple Test Corrections Types

- Family-Wise Error Rate (FWER): Controls the probability that any test is a false positive
 - Bonferroni Correction: Very stringent correction
 - Adjusted p-value = original p-value * number of tests
- False Discovery Rate (FDR): Controls the proportion of tests that are false positives
 - Widely used alternative to FWER (e.g. Bonferroni correction)
 - Example (next slide)

Benjamini-Hochberg (BH) FDR Procedure

- Goal: Calculate the new p-value cutoff for a given set of p-values
- Assume $n = 100, \alpha = 0.05$
 - α denotes the desired false positive rate,

1. Rank n p-values (large to small)

2. Calculate q-values

$$q = \alpha \times \frac{n - \text{rank} + 1}{n}$$

3. Select the lowest ranked p-value that is lower than α

Benjamini-Hochberg Example

p-value	Rank	q-value	p < q
0.9	1	$0.05 \cdot (100 - 1 + 1) / 100 = 0.05$	FALSE
0.7	2	$0.05 \cdot (100 - 2 + 1) / 100 = 0.0495$	FALSE
0.5	3	$0.05 \cdot (100 - 3 + 1) / 100 = 0.049$	FALSE
0.04	4	$0.05 \cdot (100 - 4 + 1) / 100 = 0.0485$	TRUE
...
0.005	n	$0.05 \cdot (100 - n + 1) / 100 = 5E-4$	FALSE

Bonferroni and Benjamini-Hochberg Corrections in R

```
pVals <- read.table("files/pvalsExample.txt")  
head(pVals$V1, 5)
```

```
[1] 0.0001264 0.0001150 0.0000113 0.0000882 0.0000190
```

```
pValsAdjusted <- p.adjust(pVals$V1, method="bonferroni")  
head(pValsAdjusted, 5)
```

```
[1] 0.0039184 0.0035650 0.0003503 0.0027342 0.0005890
```

```
# 'fdr' or 'BH' for Benjamini-Hochberg method  
pValsAdjusted <- p.adjust(pVals$V1, method="fdr")  
head(pValsAdjusted, 5)
```

```
[1] 0.00013061333 0.00012293103 0.00002097059  
0.00010126667 0.00002804762
```

Additional Enrichment Analyses

- Gene Set Enrichment Analysis (GSEA): GSEA is one of the best known enrichment analyses
 - This method additionally takes into account numeric values associated with the genes (e.g. gene expression levels)
 - They provide many collections of “gene sets” that can be used with GSEA or related methods
 - <http://software.broadinstitute.org/gsea/msigdb>

Regression

- Goal: Find the relationship between an independent variable and a set of dependent variables (also known as predictor or features)
 - Example: The relationship between the expression some genes and drug response

Given n observations each with a response variable y and p predictors (or features)

$$Y = (y_1, \dots, y_n)^T, \quad n \times 1$$
$$X = (X_1, \dots, X_p), \quad n \times p$$

Goal: We want to find a set of regression coefficients β for $x = (x_1, \dots, x_p)$ to describe the relationship between y and x_1, \dots, x_p

$$\hat{y} = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots$$

- \hat{y} is the predicted value
- $\hat{\beta}$ are the estimated regression coefficients (as opposed to the true coefficients)

Example Regression

```
fit <- lm(Petal.Width ~ Petal.Length, data=iris)
summary(fit)
```

Call:

```
lm(formula = Petal.Width ~ Petal.Length, data = iris)
```

Residuals:

Min	1Q	Median	3Q	Max
-0.56515	-0.12358	-0.01898	0.13288	0.64272

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-0.363076	0.039762	-9.131	4.7e-16	***
Petal.Length	0.415755	0.009582	43.387	< 2e-16	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.2065 on 148 degrees of freedom

Multiple R-squared: 0.9271, Adjusted R-squared: 0.9266

F-statistic: 1882 on 1 and 148 DF, p-value: < 2.2e-16

Interpreting lm() Results Summary

- Residuals: The difference between the actual and predicted values
- Estimate: Regression coefficient estimates
- Std. Error: Measurement of the variability of the coefficient estimate. Lower is better.
- t value: Coefficient score to describe the importance of predictor; used to calculate the p-value
- Pr(>|t|): Coefficient p-value. Probability the predictor is **NOT** relevant
- R-square: Score for evaluating how well the model fits the data. Higher is better. This can be adjusted for the number of predictors used in the model.
 - Values ~0.7 are of more interest, but there is no standard rule
- More information: <http://blog.yhat.com/posts/r-lm-summary.html>

Some Issues with Regression

- Missing Data: What if your data has missing values?
 - Imputation can be used to fill in missing values using other data points
- Too many predictors versus the number of samples
 - The “Curse of Dimensionality” (later slide)
 - Regularized regression methods can be used to select features to be included in the model
- Overfitting: Will your model work on other datasets?
 - Excessively complex models (e.g. having too many predictors) can have poor predictive performance when tested on new data
 - Regularized regression methods have properties to avoid overfitting

Missing Data

- By default, `lm()` removes rows that contain missing values
- An alternative is imputation to fill in missing values
 - `impute` imputes using K-Nearest Neighbors (KNN)
 - Step 1: Identify K number of neighbors based on Euclidean distance
 - Step 2: Average the values of the neighbors and replace the missing value

Basic Rules for Imputation

- Should be done when the number of missing values is small
 - A safe maximum threshold is 5% of the total for large datasets
- Should be done when the imputed values are plausible for the missing values
- Should be done when it is assumed that the missing values occur at random
 - If missing values do not occur at random, the data collection should be investigated and/or the values should be dropped

The "Curse of Dimensionality"

- If the number of predictors is greater than the number of samples, it will not be possible to estimate relevant regression parameters in the full model.
 - This is due to the degrees of freedom in the system.
 - There are n observations and $p + 1$ parameters (one regression coefficient for each predictor plus the intercept) leaving $n - p - 1$ degrees of freedom.
 - Increasing the sample size provides more information about the population test.
 - Increasing the number of predictors in the resulting model lowers the degrees of freedom available to estimate the variability of the predictors; this increases the variance of the regression coefficient estimates and reduces confidence in the model.

LASSO, Ridge, and Elastic Net Regularized Regression

- Least Absolute Shrinkage and Selection Operator (LASSO): Tends to produce sparse (i.e. few predictors) whereby the algorithm selects an arbitrary predictor among a set of correlated ones
- Ridge: Tends to select all correlated predictors with their coefficient values equal to each other
- Elastic Net (EN): Blends the concepts of LASSO and Ridge regression to attempt to create a model that is both sparse, but also includes correlated predictors
 - EN parameter $\alpha = 1$ represents LASSO regression, while $\alpha = 0$ approaches Ridge regression
- LASSO, Ridge, and Elastic Net regression are available from the `glmnet` R package
 - Sousa FG et al. DNA damage response alterations and its relation with drug

Clustering

- Goal: Divide data into groups (clusters), so that group members are more “similar” to each other than to members outside the group
 - Example: Cluster a set of drugs. For drugs without a known mechanism of action (MOA), predict a potential MOA based on how the unknown MOA drugs cluster with known MOA drugs
- Clustering differs from classification in that in classification we have known groups


Hierarchical Clustering

- R uses an agglomerative (bottom-up) clustering approach
 - Alternative: Divise (top-down) is similar to agglomerative, but in reverse
- Algorithm
 1. All points start in their own clusters
 2. At each iteration merge the 2 most similar structures
 3. Stop if there is a single cluster containing all points, else go to Step 2

Hierarchical Clustering Example

- Heatmap shows the expression of 20 oncogenes from 20 NCI-60 celllines

```
dat <- read.table("files/heatmapExample.txt", sep="\t", header=TRUE)  
mat <- as.matrix(dat); heatmap(mat, cexCol=0.75)
```



Cluster Similarity (Linkage)

- Distances between clusters are calculated to determine cluster similarity
- Options of Cluster Linkage Distances
 - Single: Distance between two clusters is defined by the distance between the two closest points.
 - Average: Average of all pairwise distances between the points in two clusters
 - Complete: Distance between two clusters is defined by the distance between the two farthest points.
- `hclust()` used by `heatmap()` in R uses the “complete” method by default

K-Means Clustering


- Algorithm
 1. A user-selected (k) number of means are randomly generated from the data
 2. k clusters are created by grouping data points to the nearest mean.
 3. The centroid of the clusters becomes the new mean.
 4. Steps 2 and 3 are repeated until the clusters do not change anymore

K-Means Clustering Example

```
# Retain only the numeric data in the iris dataset
iris_data <- iris[, 1:4]

# nstart: try multiple initial configurations and report the best one
kc <- kmeans(iris_data, 3, nstart=25)

par(mai=c(1,1,0,0))
plot(iris[,c("Sepal.Length", "Sepal.Width")], bg=c("red","green3","blue")[kc$cluster],
     pch=21)
points(kc$centers[,c("Sepal.Length", "Sepal.Width")], col=c("red","green3","blue"),
       pch=8, cex=2)
```




Determine K-Means Cluster Quality

```
library(cluster)
dataDist <- dist(iris_data)
si <- silhouette(kc$cl,
dataDist)
```

```
plot(si, col = c("red",
"green3", "blue"))
```

- Silhouette Plot
 - Horizontal barplot is the goodness of fit of sample within the cluster
 - Longer is better
 - Rightmost number, S_i , is average length
- Average Silhouette Guidelines
 - 0.71-1.0: Strong clustering
 - 0.51-0.70: Reasonable clustering
 - < 0.50 : Weak clustering




Selecting k with Average Silhouette

```
library(cluster)
kMax <- 15
avgSi <- rep(0, kMax)

# Average silhouette width
# k: 2 to 15
for(i in 2:kMax){
  results <-
kmeans(iris_data,
centers=i)
  si <-
silhouette(results$cluster,
dist(iris_data))
  avgSi[i] <- mean(si[,
"sil_width"])
}
```

```
plot(1:kMax, avgSi,
type="b", pch=19,
xlab="Number of Clusters
(k)")
abline(v=which.max(avgSi),
lty=2)
```



Compare Known Classes with Clusters

```
table(iris$Species, kc$cluster)
```

	1	2	3
setosa	50	0	0
versicolor	0	48	2
virginica	0	14	36

Differences between Hierarchical and K-Means Clustering

- Clusters
 - K-means produces a single set of clusters
 - Hierarchical produces different clusters depending on where the tree is cut
- Cluster Number
 - K-means requires the number clusters to be set
 - Hierarchical clustering does not require the number of clusters to be set
- Speed
 - K-means is faster than hierarchical clustering

Dimensionality Reduction

- Goal: Seeks to reveal correlations that exist in data with many predictors
 - This can be used to reduce the dimensionality of the data by decreasing the redundancy of correlated variables
- Principal Component Analysis is a method for dimension reduction

What is Principal Component Analysis (PCA)

- Goal: PCA seeks to simplify a multi-dimensional (e.g. one with 3+ predictors (features)) dataset
 - Used for feature extraction
 - May reveal clusters and help validate clustering results
- Example: We have a dataset with 20 dimensions and it may be interesting to plot the data in two dimensions
- Results:
 - Loadings: Weights for the original values to get the component scores
 - Component Scores: Transformed values for a given point
- `prcomp` and `princomp` can do PCA in R; `prcomp` is the advised function

What are Principal Components?

- Each “principal component” (PC) is an axis that captures the most variance
 - Variance is a measure of the spread of data points; standard deviation is the square root of variance
 - Each PC is a combination of the original variables scaled by a coefficient
 - Every PC explains some variance
 - Each additional PC explains less variance than the previous one

PCA Example

- Using scale=TRUE is advisable
- prcomp first transforms the data by centering and scaling
 - Centering is done by subtracting the column means
 - Scaling is done by dividing the (centered) columns of X by their standard deviations

```
iris_data <- iris[, 1:4]


pcaResult <- prcomp(iris_data, scale=TRUE)
summary(pcaResult)
```

Importance of components:


	PC1	PC2	PC3	PC4
Standard deviation	1.7084	0.9560	0.38309	0.14393
Proportion of Variance	0.7296	0.2285	0.03669	0.00518
Cumulative Proportion	0.7296	0.9581	0.99482	1.00000

PCA Example Plots

```
# First 2 principal components (PC)
plot(pcaResult$x, pch=21,
     bg=c("red", "green3", "blue"),
     [unclass(iris$Species)])
```



```
# PC variances
plot(pcaResult,
     type="line", cex.lab=1.5,
     cex.main=1.5, main="")
abline(h=1, lty=3,
       col="red")
```



Recovering the Original Data

```
# Weights (known as loadings)
pcaResult$rotation
```

	PC1	PC2	PC3	PC4
Sepal.Length	0.5210659	-0.37741762	0.7195664	0.2612863
Sepal.Width	-0.2693474	-0.92329566	-0.2443818	-0.1235096
Petal.Length	0.5804131	-0.02449161	-0.1421264	-0.8014492
Petal.Width	0.5648565	-0.06694199	-0.6342727	0.5235971

```
# Original
iris_data[1,]
```

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width
1	5.1	3.5	1.4	0.2

```
# Transformed
pcaResult$x[1,]
```

	PC1	PC2	PC3	PC4
	-2.25714118	-0.47842383	0.12727962	0.02408751


```
# Recovered
tmp <- t(t(pcaResult$x %*% t(pcaResult$rotation)) * pcaResult$scale + pcaResult$center)
tmp[1,]
```

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width
	5.1	3.5	1.4	0.2

Visualizing PCA Results with Biplots

- Visualizes the magnitude and sign of each feature's contribution to a PC
- Visualizes each observation in terms of PCs
- Closeness equals similarity for points and vectors

```
biplot(pcaResult, scale=0, cex=.7)
```



Correlations between Vectors

Feature to Principal Component Correlations

```
cor(iris_data, pcaResult$x)
```

	PC1	PC2	PC3
PC4			
Sepal.Length	0.8901688	-0.36082989	0.27565767
0.03760602			
Sepal.Width	-0.4601427	-0.88271627	-0.09361987
-0.01777631			
Petal.Length	0.9915552	-0.02341519	-0.05444699
-0.11534978			
Petal.Width	0.9649790	-0.06399985	-0.24298265
0.07535950			

Feature to Principal Component Contributions

```
tmp <- abs(pcaResult$rotation)
sweep(tmp, 2, colSums(tmp), "/" )
```

	PC1	PC2	PC3	PC4
Sepal.Length	0.2691897	0.27110474	0.41346137	0.15281309
Sepal.Width	0.1391485	0.66321714	0.14042128	0.07223451

How Many Principal Components Should be Kept?

- Kaiser criterion
 - Retain only principal components that with a variance greater than 1
 - Simple, but less advisable
- Scree Test
 - Find the place where the smooth decrease in the variances levels off
 - Multiple users may interpret the data different, unless trained the same

Getting Help

- Cross-Validated Stats Exchange
 - Part of Stack Overflow
 - <http://stats.stackexchange.com/>
- Biostars
 - <https://www.biostars.org>